

COGNIZANT TECHNOVERSE HACKATHON 2026 • NATIONAL FINALS

AUTHREX

The auditable AI copilot for oncology prior authorization.

Verify clinical. Cite policy. Decide in 60 seconds.

Team AeroFyta

Danish A. G. (Lead) · Preethi Sivachandran · Sanjay N · Gayathri B

Chennai Institute of Technology · 2027 Engineering · Pune · May 7, 2026

1 • THE PROBLEM

Every 2 seconds.

a prior authorization is filed in the United States.

A doctor waits. A patient delays. A diagnosis ages.

14 days to a verdict. **\$1,500** per case in admin labour. **600M** PAs/year.

1 • PROBLEM STATEMENT

The cost of broken prior authorization

Four numbers that explain why CMS-0057-F mandated this be solved by January 2027.

\$30B

annual US prior-auth admin waste

Source: CAQH 2024

94%

of physicians say PA delays patient care

Source: AMA 2024

80.7%

of appealed Medicare-Advantage denials are overturned

Source: KFF 2024

38%

YoY growth in oncology PA volume

Source: ASCO QOPI 2024

2 • SOLUTION DESCRIPTION

Authrex turns 14-day fax-fights into 60-second verdicts.

A 7-agent LangGraph DAG reads FHIR clinical data + payer policy, produces a citation-grounded verdict, and auto-drafts an NCCN-cited appeal letter when the case is denied — every claim bound to a stable pointer, every byte verifiable.

7 Parent agents LangGraph DAG	22 Sub-agents bounded competencies	60s End-to-end decision verified live	\$1.01 Real per-case cost vs \$1,500 manual	8/8 CMS-0057-F clauses \$ IV mapped
--	---	--	--	--

1. PROVIDER-SIDE • 2. FHIR-NATIVE • 3. CITATION-GROUNDED AUDIT-FIRST • 4. APPEALS AS SIDE-EFFECT • 5.

3 • UNIQUENESS / INNOVATIVENESS

What no one else does — on a multi-agent DAG.

01

Provider-side, by design

We work for the doctor, not the payer. Our incentive: overturning denials, not creating them.

02

Citation-grounded by architecture

Every claim binds to a FHIR resource ID + policy section. Cannot ship without citations.

03

Appeals as side-effect

When we deny, the NCCN-cited appeal letter is already drafted. Doctor never starts from blank.

04

Hash-chained audit ledger

SHA-256 chain in Postgres triggers. Tamper-evident by kernel constraint, not by hope.

05

AMBIGUOUS as 1st-class verdict

Three outcomes (APPROVE/REFER/DENY). Refuse to manufacture verdicts when evidence is missing.

06

Cognizant-stack native

Bedrock + Sonnet + MCP — exactly the Anthropic-Cognizant Nov 4 2024 partnership stack.

3 · INNOVATIONS WE BROUGHT

Nine engineering decisions other teams will not make.



Pre-LLM PHI sanitisation

PHI never enters context. Structurally impossible to leak.



Deterministic verdict synthesis

LLM produces evidence; pure Python produces verdict.



Per-criterion parallel fan-out

8 criteria checked in parallel — 8× faster, more accurate.



Provider abstraction

3 LLM providers wired. Switch via env var. No lock-in.



Hash chain in DB triggers

Audit integrity enforced by Postgres, not by hope.



Reflection grading on Haiku

Every output graded for ~\$0.005, +20% accuracy gain.



AMBIGUOUS as 1st-class verdict

Three outcomes. Refuse to fake confidence.



Pre-drafted appeal letters

NCCN-cited appeal as DAG side-effect, not afterthought.



Live-toggle demo paths

Tweaks panel switches APPROVE/REFER/DENY on stage.


5 · TECHNICAL DESIGN & ARCHITECTURE

Six concentric layers · single responsibility per layer.

L6 Ops Plane	Runbooks · checklists · evidence packs · ROI calc · compliance scorecard
L5 Surface Plane	React + Vite + TS · standalone HTML showcase · SSE client · Tweaks panel
L4 Gateway Plane	LLMClient interface · Bedrock + Anthropic + OpenRouter · cost-aware router
L3 Runtime Plane	LangGraph DAG · FastAPI + uvicorn · SSE stream · case_runner worker · reflection
L2 Agent Plane	7 parents · 22 sub-agents · Pydantic v2 schemas · .txt prompts · 5-field GraderScore
L1 Data Plane	Postgres · cases · agent_runs · llm_invocations · audit_ledger · RLS · 10-yr retention

Each layer has ONE responsibility. Top-down dependency only. Cognizant Neuro-SAN AAOSA-aligned.

Seven parents. Twenty-two sub-agents. Bounded competencies.



5 · TECH STACK · BUILT ON THE COGNIZANT STANDARD

Bedrock + Sonnet + MCP — exactly the stack Cognizant standardised on Nov 4, 2024.

AI / LLM	Backend	Frontend	Infrastructure
● Claude Sonnet 4.6 (Bedrock) ● Claude Haiku 4.5 (grader) ● LangGraph (DAG orchestration) ● MCP (tool protocol) ● Anthropic + OpenRouter (fallback)	● Python 3.11+ async ● FastAPI + uvicorn ● Pydantic v2 schemas ● AsyncPG · Postgres driver ● pytest + ruff + mypy strict	● React 18 + Vite 5 ● TypeScript strict ● Tailwind CSS ● lucide-react icons ● SSE live-trace client	● AWS Bedrock (ap-south-1) ● AWS Q Business (RAG) ● ECS Fargate · RDS · S3 ● KMS + IAM + Secrets Mgr ● CloudWatch observability

No new tech for Cognizant to adopt. We use theirs.

4 • BUSINESS IMPACT

From \$1,500 manual labour to 25 cents of AI compute.

MANUAL PRIOR AUTH

\$1,500

per case • 14 days • 3 hours of physician labour

AUTHREX (verified May 5 2026)

\$0.25

per clean APPROVE • 60 seconds • 98 LLM calls

NET SAVINGS PER CASE

\$1,499.75

99.93% cost reduction

AT ONE 200-BED CANCER CENTRE: 150 oncology PAs/month × \$1,499.75 saved =


\$224,962 / month

Payback: under 0.2 days per case • ROI on Pilot tier subscription: 16,664%

4 • MARKET OPPORTUNITY

\$30B → \$200M → \$20M.


Total US prior-auth waste, addressable specialty PA market, and 5-year capture target.



4 • **BUSINESS IMPACT (continued):** Distribution path = **Cognizant Health Sciences** sells to 47 of top 50 US payers + 200+ providers. We slot into the existing procurement vehicle.

Same architecture, six markets — fixtures-only swap.

The 7-agent DAG is payload-agnostic. Each adjacent market is a prompts + fixtures swap, not a re-architecture.



Total adjacent TAM = ~6× the prior-auth opportunity. Same architecture. Same Cognizant distribution.

COGNIZANT FIT · WHY US, NOT SOMEONE ELSE

The Anthropic-Cognizant stack is our default, not our adaptation.

Anthropic stack · Nov 4 2024	Bedrock + Claude Sonnet 4.6 + MCP — exactly the stack Cognizant standardised on.
Neuro-SAN AAOSA pattern	7-parent bounded-responsibility DAG · stateful continuity · per-tenant adaptation · observability of every hop.
Cognizant Health Sciences	Oncology PA wedge fits the named 2026 vertical focus — 38% YoY oncology PA volume growth.
Cognizant Agent Foundry	Foundry-compatible manifest · model card · evaluation harness · observability spec.
TriZetto integration	Provider→payer structured-envelope submission. Mock + real adapters both written.
Kiro IDE	Spec-driven development. .kiro/specs/ on disk. Proves Cognizant IDE story without leaving the project.

We are not asking Cognizant to adopt new technology. We use theirs.

Eight of eight federal clauses in force today.

Federal mandate enforces January 2027. We are 18 months early.

8 / 8

§ IV CLAUSES IN FORCE

Verified end-to-end
on 2026-05-05

CMS-0057-F enforcement
January 2027

✓ § IV.A Auditable decisions	✓ § IV.B Citation grounding (policy)
✓ § IV.C Citation grounding (clinical)	✓ § IV.D Structured payer-API integration
✓ § IV.E Appeal window standardisation	✓ § IV.F PHI handling (minimum necessary)
✓ § IV.G Reviewer transparency	✓ § IV.H Quarterly evidence pack

7 · ROADMAP

From hackathon win to \$5M ARR · 30 months.



TEAM AEROFYTA · 4 ROLES, NO OVERLAP

Each member owns one Cognizant evaluation criterion.



Danish A. G.
Team Lead · Chief Architect

Builds: 7-agent DAG · Bedrock + Sonnet 4.6 · MCP · LLMClient · system design
Owns Q&A on: *Tech Design & Architecture · AWS / GenAI Usage*



Preethi Sivachandran
Product & Demo Lead · UX

Builds: Live demo · UX flow · Tweaks panel · standalone showcase · pitch script
Owns Q&A on: *Quality of Demo & Presentation*



Sanjay N
Backend & Compliance Engineer

Builds: FHIR R4 · Postgres + RLS · SHA-256 audit · CMS-0057-F · HIPAA / PHI
Owns Q&A on: *MVP Functional Completeness (data plane)*



Gayathri B
Healthcare Domain & Business Analyst

Builds: \$30B problem · CAQH/KFF 2024 sourcing · ROI math · pilot pipeline · NCCN
Owns Q&A on: *Problem Clarity · Business Impact*

We delivered the working system in 24 hours. We can ship at Cognizant pace tomorrow.

THE ASK

Three asks. One closing line.

ASK 1

First-prize recognition

Validation that the architecture is industrial-grade.

ASK 2

Cognizant pilot referral

1-2 Health Sciences referral accounts to onboard.

ASK 3

Internships for the team

We want to keep building Authrex inside Cognizant.

THE CLOSING LINE

*"We are not pitching a hackathon project.
We are pitching the missing piece of the Anthropic-Cognizant healthcare stack."*

THANK YOU.

Try it now on your phone.

[QR CODE]

Generate QR encoding
<https://authrex.xzashr.com>
and paste here before
printing.

authrex.xzashr.com

13 routes · 3 demo paths (APPROVE / REFER / DENY) · zero backend · runs from anywhere

TEAM CONTACTS

Danish A. G. · agdanishr@gmail.com

Preethi Sivachandran · preethisivachandran0@gmail.com

Sanjay N · Gayathri B · Chennai Institute of Technology · +91-8903099026

SOURCES & CITATIONS · ALL CLAIMS VERIFIABLE

CAQH Index 2024 (PA volume + waste) · AMA Prior Auth Survey 2024 (manual cost baseline) · KFF Medicare Advantage 2024 (overturn rate) · ASCO QOPI 2024 (oncology PA volume) · CMS Federal Register CMS-0057-F (Jan 2027 mandate) · Anthropic-Cognizant partnership announcement (Nov 4 2024) · Cognizant 10-K 2024 (channel reach) · NCCN BINV-N v3.2026 (clinical guidelines) · AWS Bedrock pricing (May 2026) · Authrex internal verified llm_invocations table (case_8f4ad9c2, 98 calls, May 5 2026)